

# Regression Modeling for Stock Prediction

```
In [2]: import numpy as np
import pandas as pd
import yfinance as yf
```

```
In [3]: # Step 1: Download the Data

tickers = ('AAPL', 'AMZN', 'MSFT', 'QQQ', '^GSPC')
start_date = '2020-01-01'
end_date = '2024-12-31'

df = yf.download(tickers, start_date, end_date)[['Close']]
df
```

C:\Users\sahil\AppData\Local\Temp\ipykernel\_1500\2293748721.py:7: FutureWarning: YF.download() has changed argument auto\_adjust default to True  
 df = yf.download(tickers, start\_date, end\_date)[['Close']]  
 [\*\*\*\*\*100%\*\*\*\*\*] 5 of 5 completed

	Ticker	AAPL	AMZN	MSFT	QQQ	<sup>^</sup> GSPC
	Date					
2020-01-02	72.468269	94.900497	152.505707	208.848953	3257.850098	
2020-01-03	71.763733	93.748497	150.606720	206.935898	3234.850098	
2020-01-06	72.335556	95.143997	150.995972	208.269241	3246.280029	
2020-01-07	71.995361	95.343002	149.619308	208.240280	3237.179932	
2020-01-08	73.153503	94.598503	152.002457	209.805450	3253.050049	
	...	...	...	...	...	...
2024-12-23	254.120682	225.059998	432.062775	520.901855	5974.069824	
2024-12-24	257.037476	229.050003	436.112885	527.965210	6040.040039	
2024-12-26	257.853790	227.050003	434.901794	527.606506	6037.589844	
2024-12-27	254.439224	223.750000	427.377319	520.593018	5970.839844	
2024-12-30	251.064484	221.300003	421.719025	513.669189	5906.939941	

1257 rows × 5 columns

```
In [4]: # Perform Features Engineering
# Shift the value by one

df['AAPL(t-1)'] = df['AAPL'].shift(1)
df['AMZN(t-1)'] = df['AMZN'].shift(1)
df['MSFT(t-1)'] = df['MSFT'].shift(1)
```

```
df['QQQ(t-1)'] = df['QQQ'].shift(1)
df['^GSPC(t-1)'] = df['^GSPC'].shift(1)
```

```
In [5]: # Moving Average 5 (Technical Indicator)
```

```
df['AAPL_MA5'] = df['AAPL'].rolling(window=5).mean()
df['AMZN_MA5'] = df['MSFT'].rolling(window=5).mean()
df['MSFT_MA5'] = df['MSFT'].rolling(window=5).mean()
df['QQQ_MA5'] = df['QQQ'].rolling(window=5).mean()
df['^GSPC_MA5'] = df['^GSPC'].rolling(window=5).mean()
```

```
In [6]: # Set Y Variable(next day)
```

```
df['Target'] = df['AAPL'].shift(-1)
```

```
In [7]: df.dropna(inplace = True)
df
```

Out[7]:

Ticker	AAPL	AMZN	MSFT	QQQ	^GSPC	AAPL(t-1)	AMZN(t-1)
Date							
2020-01-08	73.153503	94.598503	152.002457	209.805450	3253.050049	71.995361	95.343002
2020-01-09	74.707336	95.052498	153.901443	211.583237	3274.699951	73.153503	94.598503
2020-01-10	74.876251	94.157997	153.189301	211.042114	3265.350098	74.707336	95.052498
2020-01-13	76.475906	94.565002	155.031281	213.476944	3288.129883	74.876251	94.157997
2020-01-14	75.443230	93.472000	153.939423	212.636383	3283.149902	76.475906	94.565002
...	...	...	...	...	...	...	...
2024-12-20	253.344177	224.919998	433.402924	515.875793	5930.850098	248.665329	223.289993
2024-12-23	254.120682	225.059998	432.062775	520.901855	5974.069824	253.344177	224.919998
2024-12-24	257.037476	229.050003	436.112885	527.965210	6040.040039	254.120682	225.059998
2024-12-26	257.853790	227.050003	434.901794	527.606506	6037.589844	257.037476	229.050003
2024-12-27	254.439224	223.750000	427.377319	520.593018	5970.839844	257.853790	227.050003

1252 rows × 16 columns



In [8]: # See how many and which columns are in my dataframe  
df.columns

Out[8]: Index(['AAPL', 'AMZN', 'MSFT', 'QQQ', '^GSPC', 'AAPL(t-1)', 'AMZN(t-1)', 'MSFT(t-1)', 'QQQ(t-1)', '^GSPC(t-1)', 'AAPL\_MA5', 'AMZN\_MA5', 'MSFT\_MA5', 'QQQ\_MA5', '^GSPC\_MA5', 'Target'], dtype='object', name='Ticker')

In [9]: # Checking if there is any NaN values in the columns  
print(df.isna().sum())

```
Ticker
AAPL      0
AMZN      0
MSFT      0
QQQ       0
^GSPC     0
AAPL(t-1) 0
AMZN(t-1) 0
MSFT(t-1) 0
QQQ(t-1)  0
^GSPC(t-1) 0
AAPL_MA5   0
AMZN_MA5   0
MSFT_MA5   0
QQQ_MA5    0
^GSPC_MA5  0
Target     0
dtype: int64
```

In [10]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1252 entries, 2020-01-08 to 2024-12-27
Data columns (total 16 columns):
 #   Column        Non-Null Count  Dtype  
--- 
 0   AAPL          1252 non-null   float64
 1   AMZN          1252 non-null   float64
 2   MSFT          1252 non-null   float64
 3   QQQ           1252 non-null   float64
 4   ^GSPC         1252 non-null   float64
 5   AAPL(t-1)    1252 non-null   float64
 6   AMZN(t-1)    1252 non-null   float64
 7   MSFT(t-1)    1252 non-null   float64
 8   QQQ(t-1)     1252 non-null   float64
 9   ^GSPC(t-1)   1252 non-null   float64
 10  AAPL_MA5     1252 non-null   float64
 11  AMZN_MA5     1252 non-null   float64
 12  MSFT_MA5     1252 non-null   float64
 13  QQQ_MA5      1252 non-null   float64
 14  ^GSPC_MA5    1252 non-null   float64
 15  Target        1252 non-null   float64
dtypes: float64(16)
memory usage: 166.3 KB
```

In [11]: `# Run Linear Regression Model`

```
import statsmodels.api as sm

X = df[ [ 'AAPL(t-1)', 'AMZN(t-1)', 'MSFT(t-1)', 'QQQ(t-1)', '^GSPC(t-1)', 'AAPL_MA5', 'MSFT_MA5', 'QQQ_MA5', '^GSPC_MA5' ] ]
Y = df[ 'Target' ]

x_Const = sm.add_constant(X) # Intercept

# Train the model
```

```
model = sm.OLS(Y, x_Const).fit()  
  
# Summary of the Model  
model.summary()
```

Out[11]:

## OLS Regression Results

<b>Dep. Variable:</b>	Target	<b>R-squared:</b>	0.993
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.992
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1.827e+04
<b>Date:</b>	Fri, 21 Nov 2025	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	23:55:42	<b>Log-Likelihood:</b>	-3381.7
<b>No. Observations:</b>	1252	<b>AIC:</b>	6783.
<b>Df Residuals:</b>	1242	<b>BIC:</b>	6835.
<b>Df Model:</b>	9		
<b>Covariance Type:</b>	nonrobust		
	<b>coef</b>	<b>std err</b>	<b>t</b> <b>P&gt; t </b> <b>[0.025</b> <b>0.975]</b>
<b>const</b>	0.4174	1.138	0.367 0.714 -1.815 2.650
<b>AAPL(t-1)</b>	0.4601	0.081	5.703 0.000 0.302 0.618
<b>AMZN(t-1)</b>	0.0068	0.010	0.695 0.487 -0.012 0.026
<b>MSFT(t-1)</b>	-0.0132	0.052	-0.253 0.800 -0.116 0.089
<b>QQQ(t-1)</b>	0.0499	0.090	0.556 0.578 -0.126 0.226
<b>^GSPC(t-1)</b>	0.0040	0.007	0.574 0.566 -0.010 0.018
<b>AAPL_MA5</b>	0.5243	0.082	6.430 0.000 0.364 0.684
<b>AMZN_MA5</b>	0.0129	0.026	0.487 0.626 -0.039 0.065
<b>MSFT_MA5</b>	0.0129	0.026	0.487 0.626 -0.039 0.065
<b>QQQ_MA5</b>	-0.0530	0.091	-0.583 0.560 -0.231 0.125
<b>^GSPC_MA5</b>	-0.0043	0.007	-0.610 0.542 -0.018 0.010
<b>Omnibus:</b>	26.686	<b>Durbin-Watson:</b>	0.805
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	51.277
<b>Skew:</b>	-0.089	<b>Prob(JB):</b>	7.33e-12
<b>Kurtosis:</b>	3.975	<b>Cond. No.</b>	7.69e+16

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 8.03e-24. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [12]: # Run Linear Regression Model

import statsmodels.api as sm

X = df[ [ 'AAPL(t-1)', '^GSPC(t-1)' ] ]

Y = df[ 'Target' ]

x_Const = sm.add_constant(X) # Intercept

# Train the model
model = sm.OLS(Y, x_Const).fit()

# Summary of the Model
model.summary()
```

Out[12]:

## OLS Regression Results

<b>Dep. Variable:</b>	Target	<b>R-squared:</b>	0.992			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.992			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	7.762e+04			
<b>Date:</b>	Fri, 21 Nov 2025	<b>Prob (F-statistic):</b>	0.00			
<b>Time:</b>	23:55:42	<b>Log-Likelihood:</b>	-3420.7			
<b>No. Observations:</b>	1252	<b>AIC:</b>	6847.			
<b>Df Residuals:</b>	1249	<b>BIC:</b>	6863.			
<b>Df Model:</b>	2					
<b>Covariance Type:</b>	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
<b>const</b>	-0.7172	0.748	-0.959	0.338	-2.185	0.750
<b>AAPL(t-1)</b>	0.9840	0.007	140.621	0.000	0.970	0.998
<b>^GSPC(t-1)</b>	0.0008	0.000	2.114	0.035	5.8e-05	0.002
<b>Omnibus:</b>	40.865	<b>Durbin-Watson:</b>	1.041			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	103.864			
<b>Skew:</b>	-0.036	<b>Prob(JB):</b>	2.79e-23			
<b>Kurtosis:</b>	4.409	<b>Cond. No.</b>	3.08e+04			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.08e+04. This might indicate that there are strong multicollinearity or other numerical problems.

In [13]: # Predict the value using the Actual Price

```
df_train_predict = pd.DataFrame()
df_train_predict['Actual'] = df['Target']
df_train_predict['Predicted'] = model.predict(x_Const)
df_train_predict
```

Out[13]:

	Actual	Predicted
Date		
2020-01-08	74.707336	72.734141
2020-01-09	74.876251	73.886544
2020-01-10	76.475906	75.432966
2020-01-13	75.443230	75.591649
2020-01-14	75.119904	77.184070
...	...	...
2024-12-20	254.120682	248.696998
2024-12-23	257.037476	253.352384
2024-12-24	257.853790	254.151282
2024-12-26	254.439224	257.074565
2024-12-27	251.064484	257.875851


1252 rows × 2 columns

In [14]:

```
# Plot between Actual vs Predict

import matplotlib.pyplot as plt

plt.figure( figsize = (14,5))
plt.plot(df_train_predict.index, df_train_predict['Actual'], label='Actual', color='red')
plt.plot(df_train_predict.index, df_train_predict['Predicted'], label='Predicted', color='blue')
plt.title('Plot between Actual v/s Predicted Price')
plt.xlabel('Date')
plt.ylabel('Price')
plt.show()
```




# Check the assumption of the LR model

```
In [16]: # 1) Linearity between dependent and independent

import seaborn as sns

df = df[['AAPL', 'AAPL(t-1)', '^GSPC(t-1)', 'AAPL_MA5']]
sns.pairplot(df)
```

Out[16]: <seaborn.axisgrid.PairGrid at 0x12b425a0d70>



```
In [17]: # 2) Multicollinearity => using VIF
```

```
from statsmodels.stats.outliers_influence import variance_inflation_factor

vif = pd.DataFrame()
vif['features'] = x_Const.columns
```

```
vif['vif'] = [variance_inflation_factor(x_Const.values, i) for i in range(x_Const.s
vif = vif[1:]
vif
```

Out[17]:

	features	vif
1	AAPL(t-1)	7.634911
2	^GSPC(t-1)	7.634911


In [18]:

```
# Homoscedasticity => fitting residual with predicted value

residual = model.resid # Actual - Predicted
fitted = model.fittedvalues # Predict Y Value

plt.figure( figsize = (7,5))
sns.scatterplot(x = fitted, y = residual)
plt.axhline(0, color='black', linestyle='--')
plt.title('Testing Homoscedasticity')
plt.xlabel('Fitted')
plt.ylabel('Residual')
plt.show()


# This graph does not Look Like the funnel
```



In [19]:


```
# 4) Normality of residual => Histogram or QQ plot
# Histogram
```

```
plt.figure(figsize = (8,6))
plt.hist(residual, bins=30)
plt.title('Histogram for Normality')
plt.show()
```



In [20]: # QQ Plot

```
sm.qqplot(residual, line ='45', fit = True)
plt.title('QQ PPlot')
plt.show()
```



```
In [21]: # Auto Correlation of Residual => Durbin Watson Test
from statsmodels.stats.stattools import durbin_watson
dw = durbin_watson(residual)
```

```
Out[21]: 1.0410072213200119
```

## Predict the stock price for the year 2025

```
In [23]: # Step 1: Download the data from Yahoo Finance
tickers = ('AAPL', '^GSPC')
start_date = '2025-01-01'
end_date = '2025-03-31'

df = yf.download(tickers, start_date, end_date)[['Close']]
df.head()
```

```
C:\Users\sahil\AppData\Local\Temp\ipykernel_1500\3989341984.py:7: FutureWarning: YF.
download() has changed argument auto_adjust default to True
  df = yf.download(tickers, start_date, end_date)[['Close']]
[*****100%*****] 2 of 2 completed
```

Out[23]:

Ticker	AAPL	$^GSPC$
Date		
2025-01-02	242.752090	5868.549805
2025-01-03	242.264297	5942.470215
2025-01-06	243.896912	5975.379883
2025-01-07	241.119476	5909.029785
2025-01-08	241.607269	5918.250000

In [24]: # Step 2: Feature Engineer

```
df['AAPL(t-1)'] = df['AAPL'].shift(1)
df['^GSPC(t-1)'] = df['^GSPC'].shift(1)
df = df.dropna()
df.head()
```

Out[24]:

Ticker	AAPL	$^GSPC$	$\text{AAPL(t-1)}$	$\text{^GSPC(t-1)}$
Date				
2025-01-03	242.264297	5942.470215	242.752090	5868.549805
2025-01-06	243.896912	5975.379883	242.264297	5942.470215
2025-01-07	241.119476	5909.029785	243.896912	5975.379883
2025-01-08	241.607269	5918.250000	241.119476	5909.029785
2025-01-10	235.783600	5827.040039	241.607269	5918.250000

In [25]:

```
x_test = df[ ['AAPL(t-1)', '^GSPC(t-1)'] ]
x_test = sm.add_constant(x_test)


df_result = pd.DataFrame()
df_result['Actual'] = df['AAPL']
df_result['Predicted'] = model.predict(x_test)
df_result.head()
```

Out[25]:

	Actual	Predicted
Date		
2025-01-03	242.264297	242.879509
2025-01-06	243.896912	242.459055
2025-01-07	241.119476	244.092067
2025-01-08	241.607269	241.305608
2025-01-10	235.783600	241.793027

In [26]: # Plot Actual V/S Predicted

```
plt.figure( figsize = (14,7))
plt.plot(df_result.index, df_result['Actual'], label='Actual', color='red')
plt.plot(df_result.index, df_result['Predicted'], label='predict', color='black')
plt.title('Plot the graph Actual v/s Predict')
plt.xlabel('Date')
plt.ylabel('Stock Price')
plt.grid(True)
plt.show()
```



In [27]: # Risk Metrics => Calculate RMSE and MSE

```
from sklearn.metrics import mean_squared_error, r2_score

r2 = r2_score(df_result['Actual'], df_result['Predicted'])

mse = mean_squared_error(df_result['Actual'], df_result['Predicted']) # MSE

rmse = np.sqrt(mse) # RMSE

print('R2 =', r2)
print("MSE:", mse)
print("RMSE:", rmse)
```

R2 = 0.8348121724833452  
MSE: 17.23421034626402  
RMSE: 4.151410645342619

In [ ]: