

## 1. Introduction

```
In [30]: # This notebook compares four regression models – Linear, Lasso, Ridge, and Elastic
# from Yahoo Finance. The aim is to analyze performance, interpret results, and ide
# Mention data source (Yahoo Finance) and time frame (2020-2024).
```

## 2. Data Overview

```
In [4]: import pandas as pd
import yfinance as yf
import numpy as np

tickers = ('AAPL', 'AMZN', 'MSFT', 'QQQ', '^GSPC')
start_date = '2020-01-01'
end_date = '2024-12-31'

df = yf.download(tickers, start_date, end_date)[['Close']]
df.head()
```

C:\Users\sahil\AppData\Local\Temp\ipykernel\_26932\1826400475.py:9: FutureWarning: Y  
F.download() has changed argument auto\_adjust default to True  
df = yf.download(tickers, start\_date, end\_date)[['Close']]  
[\*\*\*\*\*100%\*\*\*\*\*] 5 of 5 completed

	<b>Ticker</b>	<b>AAPL</b>	<b>AMZN</b>	<b>MSFT</b>	<b>QQQ</b>	<b>^GSPC</b>
	<b>Date</b>					
<b>2020-01-02</b>	72.468269	94.900497	152.505676	208.848953	3257.850098	
<b>2020-01-03</b>	71.763733	93.748497	150.606766	206.935944	3234.850098	
<b>2020-01-06</b>	72.335579	95.143997	150.996033	208.269257	3246.280029	
<b>2020-01-07</b>	71.995331	95.343002	149.619278	208.240295	3237.179932	
<b>2020-01-08</b>	73.153465	94.598503	152.002472	209.805466	3253.050049	

```
In [32]: df.describe()
```

Out[32]:

Ticker	AAPL	AMZN	MSFT	QQQ	^GSPC
<b>count</b>	1257.000000	1257.000000	1257.000000	1257.000000	1257.000000
<b>mean</b>	151.581359	146.655994	286.891954	338.440789	4258.316540
<b>std</b>	41.855809	31.951644	81.367507	78.809634	766.387561
<b>min</b>	54.264336	81.820000	128.929977	163.573898	2237.399902
<b>25%</b>	126.484459	120.970001	227.434586	283.099030	3818.830078
<b>50%</b>	150.220779	153.839996	274.704865	327.713470	4204.310059
<b>75%</b>	176.073013	170.000000	331.681458	380.650818	4602.450195
<b>max</b>	257.853760	232.929993	462.375580	535.281128	6090.270020

### 3. Model Summary

In [ ]:

Model	Regularization Type	Key Parameters	R <sup>2</sup>	MSE	RMSE
Linear	None	-	0.83	17.23	4.15
Lasso	L1	alpha=0.1	0.75	34.32	5.85
Ridge	L2	alpha=1.0	0.74	35.98	5.99
ElasticNet	L1 + L2	alpha=0.1, l1_ratio=0.5	0.75	35.00	5.91

### 4. Visualization of RMSE Comparison between Model

In [23]:

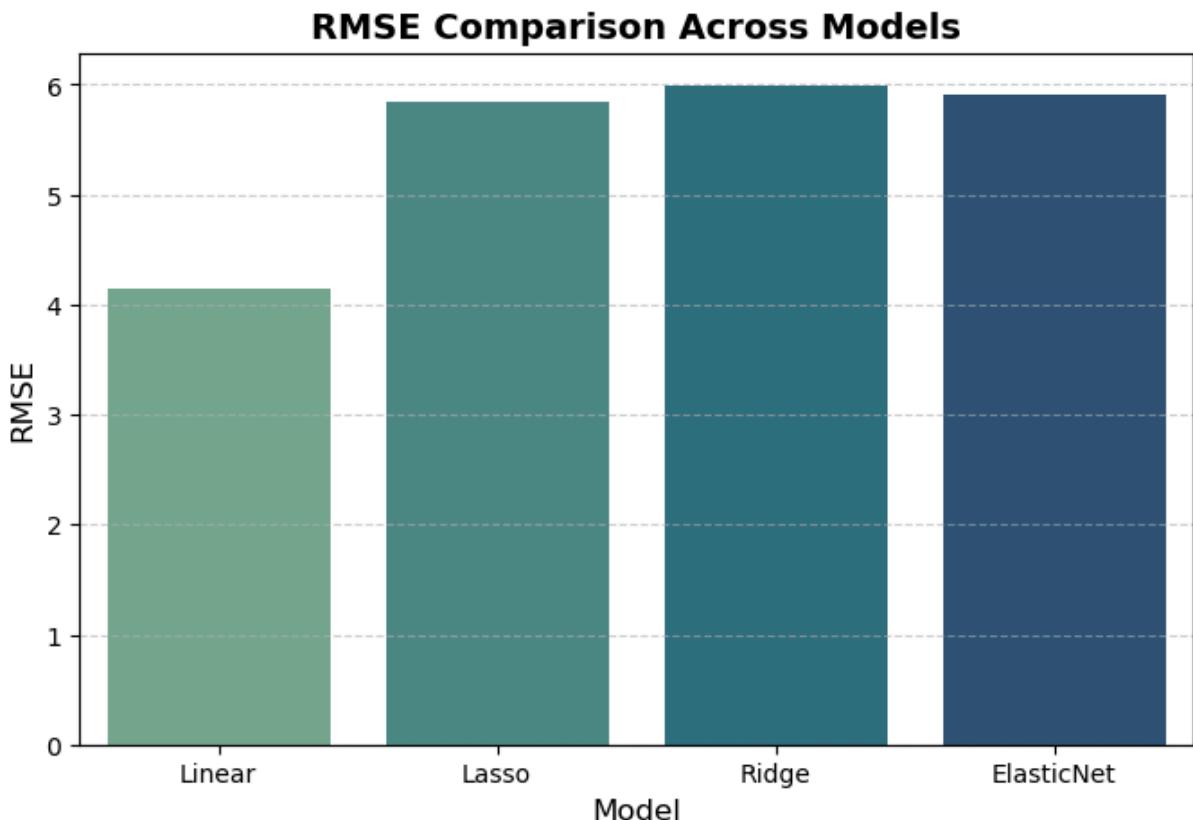
```

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Create the model summary DataFrame
model_summary = pd.DataFrame({
    'Model': ['Linear', 'Lasso', 'Ridge', 'ElasticNet'],
    'Regularization Type': ['None', 'L1', 'L2', 'L1 + L2'],
    'RMSE': [4.15, 5.85, 5.99, 5.91]
})
# Plot RMSE comparison
plt.figure(figsize=(8, 5))
sns.barplot(x='Model', y='RMSE', data=model_summary, palette='crest')
plt.title('RMSE Comparison Across Models', fontsize=14, fontweight='bold')
plt.xlabel('Model', fontsize=12)
plt.ylabel('RMSE', fontsize=12)
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.show()

```

```
C:\Users\sahil\AppData\Local\Temp\ipykernel_26932\1419189477.py:13: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.1  
4.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.  
  
sns.barplot(x='Model', y='RMSE', data=model_summary, palette='crest')
```



```
In [ ]: # Observation of the graph:  
# The RMSE comparison shows that Linear Regression achieved the Lowest error (~ 4.1  
# Regularized models – Lasso, Ridge, and ElasticNet – produced slightly higher RMSE  
# for improved stability and reduced overfitting.  
# This highlights the classic trade-off between model complexity and generalization
```

## 5. Interpretation

```
In [ ]: # Which model performed best and why:  
  
# Linear Regression achieved the lowest RMSE (~ 4.15) and the highest R2 (~ 0.83),  
# However, since it lacks regularization, it may overfit and perform less reliably o  
# Ridge, Lasso, and ElasticNet introduced regularization to control overfitting.  
# Their RMSE values (~ 5.8–6.0) were slightly higher, but they produced more stable  
# Among these, ElasticNet offered the most balanced performance – combining Ridge's  
  
# Overfitting Analysis  
  
# Linear Regression: Lowest bias, highest variance – fits training data closely but  
# Ridge & ElasticNet Introduced regularization penalties, reducing variance and imp  
# Lasso Slightly higher bias but simpler and more interpretable due to coefficient s
```

```
# Bias-Variance Trade-Off:
```

Model	Bias	Variance	Generalization
Linear	Low	High	Moderate
Lasso	High	Low	Good
Ridge	Moderate	Low	Strong
ElasticNet	Balanced	Balanced	Strong

```
#Feature Importance:
```

```
# Lasso Regression highlighted key predictors such as AAPL(t-1) and AAPL_MA5, setting up a baseline for comparison.
# Ridge Regression retained all features but reduced the weight of correlated ones, providing a more stable fit.
# ElasticNet Regression combined both methods – performing partial feature selection.
```

## 6. Key Learnings

In [ ]:

- # Linear Regression
  - Served as the baseline model.
  - Provided a clear starting point for comparison and interpretation.
  - Highlighted the limitations of unregularized models (tendency to overfit).
  
- # Lasso Regression
  - Demonstrated feature selection through L1 regularization.
  - Simplified the model by removing non-essential predictors.
  - Best when interpretability and sparsity are desired.
  
- # Ridge Regression
  - Offered the most stable and reliable predictions.
  - Controlled coefficient magnitudes, reducing sensitivity to multicollinearity.
  - Showed strong generalization across the dataset.
  
- # ElasticNet Regression
  - Combined Lasso's feature selection and Ridge's stability
  - Provided a balanced performance between accuracy and interpretability.
  - Ideal when predictors are correlated or when both sparsity and stability are valuable.